

Gaurav Lamba

+91 7827472443 | gauravlamba1109@gmail.com | [linkedin.com](https://www.linkedin.com/in/gauravlamba) | [github.com](https://github.com/gauravlamba) | [leetcode.com](https://leetcode.com/gauravlamba)

EDUCATION

Netaji Subhas University of Technology

Bachelor of Technology in Electrical Engineering | CGPA: 7.89

New Delhi, India

Aug. 2019 – May 2023

WORK EXPERIENCE

Software Development Engineer

June 2023 – Present

Oracle

Bangalore, India

- Owned design and delivery of a **Direct-to-Management-Ledger** configuration capability using Java, Jakarta REST, React, TypeScript, Oracle SQL, and Liquibase; added RBAC, validation, versioned audit history, and downstream workflow integration.
- Engineered a **SQL execution-tracking system** for PMF resume operations with per-statement state, idempotent resume, and granular retries, reducing resume time by **75%** and avoiding full delete-and-rerun cycles.
- Extended daily and period data compaction with current-day and rolling seven-day modes, then optimized high-volume Oracle paths with partition drops, CTAS cleanup, and configurable optimizer hints; reduced a critical cleanup path from **30s to <2s**.
- Added processing-run identifiers to Fact Accounting Entries and incremental Revaluation results-area flows using Java, Oracle SQL, and Liquibase, improving run-level traceability across high-volume accounting pipelines.
- Spearheaded migration of KYC microservices from **Spring Boot to Helidon**, reducing startup time by **40%** and memory usage by **25%**; contributed Helidon 4 and JDK 21 compatibility work across BCE services.
- Delivered Revaluation Management-Ledger defaults through a searchable React interface and transactional Java APIs with audit logging, dimension-member validation, and fiscal-period controls that prevent invalid financial processing.
- Integrated the FetchRates REST API into BCE Revaluation and hardened exchange-rate loading against intermittent responses and null-data failures, improving reliability of FX-driven accounting runs.
- Architected a full-stack performance monitoring platform using React and Spring Boot with custom REST APIs over **20+ SQL diagnostics**; adopted by **15+ engineers**, reducing analysis time from **30 minutes to 5 minutes**.
- Implemented legal-entity fiscal-period validation for BCE Day 0 and Revaluation workflows, integrating request validation, execution history, PMF dashboard events, and resumable failure handling to block processing against closed periods.
- Received the **Oracle Pacesetter Award** for improving QA productivity through the performance monitoring platform.

Product Analytics Intern

Nov. 2021 – Feb. 2022

Cympl Studios

New Delhi, India

- Analyzed gameplay data for mobile apps with **1M+ downloads** using BigQuery and Tableau, extracting KPIs to guide feature prioritization and optimizing analytical SQL latency by **40%**.

TECHNICAL SKILLS

Languages & Frameworks: Java, Helidon, Spring Boot, JavaScript, TypeScript, React.js, Node.js, Python, C++, Kotlin, SQL

Tools & Technologies: Kubernetes, Docker, Git, OCI, BigQuery, Tableau, Liquibase, Fortify, OWASP, CI/CD

Databases & Architecture: Oracle DB, PostgreSQL, Firebase, Microservices, Distributed Systems, System Design, REST APIs, JDBC, WebRTC

PROJECTS

ApnaCam | Kotlin, Firebase, MVVM | Play Store

May 2022 – Nov. 2022

- Built an Android security-camera application reaching **10,000+ downloads**; implemented MVVM with Data Binding and Coroutines and integrated YouTube Data API and REST APIs.

MCP Defect Resolution Workspace | MCP, Codex, Git, Kubernetes

2026

- Integrated Jira MCP and BugDB MCP to present a selectable defect queue and retrieve descriptions, logs, comments, and reproduction steps; used Confluence MCP for evidence tracking and validated issues against UI behavior and source code to reject false positives.
- Automated isolated branch creation, fix implementation, test-environment deployment, and PMF regression execution; rechecks resolution before requesting commit approval and records findings in Confluence.
- Reduced agent token usage by 40% using deterministic tools and leader scripts, field-selective MCP retrieval, bounded and deduplicated logs, cached tool results, and structured state summaries.